

# **CSEN603**

## **Assignment II**

Omar Sherif  
Amir George

## Question 1

### Hash Joins

We discuss three variations of the Hash Join.

- Simple Hash Join
- Grace/Partition Hash Join
- Hybrid Hash Join

We use example relations  $S$  with  $m$  tuples and  $R$  with  $n$  tuples where  $S$  is the smaller relation, and  $M$  memory buffers where each buffer can hold a relation block and  $A, B$  are  $S$  and  $R$ 's join attributes respectively.

### Simple Hash Join

**Partition Phase** Relation  $S$  is partitioned using a hash function  $h$  on the join attribute into  $M - 1$  buckets by loading  $S$ 's blocks one by one into the remaining buffer emptying and saving each bucket to secondary memory as it is filled.

**Join/Probe Phase** Having  $S$ 's  $M - 1$  buffers in memory  $R$ 's blocks are loaded one by one in the remaining buffer and each tuple  $r$  in the block is applied to  $h$  on its join attribute  $A$ , for each tuple  $s$  in bucket  $h(r.A)$  if  $s.B = r.A$  where  $B$  is the join attribute tuple  $r \circ s$  is written to the output buffer.

---

Hashing  $S$

---

```
foreach tuple  $s \in S$  do
  put  $s$  in bucket  $h(s.A)$ 
end
```

---



---

Probing

---

```
foreach tuple  $r \in R$  do
  foreach tuple  $s$  in bucket  $h(r.A)$  do
    if  $r.A = s.B$  then
      output  $r \circ s$ 
    end
  end
end
```

---

Two relations  $Guest(Name, Roomno)$ ,  $Room(Roomno, telephone)$ , we assume 4 buffers of memory are available and  $h(x) = x \bmod 3$ , where a block holds two tuples

Name	Roomno	Roomno	telephone
Osama	2	22	21849
Ahmed	9	2	21239
Yehia	55	55	2354
Bigby	47	11	211222
Wander	33	15	23521
		9	143556

The result of Hashing *Guest*

0	1	2	3
Ahmed 9 Wander 33	Yehia 55	Bigby 47 Osama 2	

after loading the first block of Room

0	1	2	3
Ahmed 9 Wander 33	Yehia 55	Bigby 47	22 21849 2 21239

22 is hashed to bucket 1 but since the only tuple in that bucket doesn't have the same value of the join attribute no joining is done, but 2 is found in bucket 2 and joined and tuple **Osama 2 21239** is outputted, then the next block of room is loaded and the process continues.

### Grace/Partition Hash Join

**Partition Phase** Relation  $S$  is partitioned using a hash function  $h$  on the join attribute into  $M - 1$  buckets by loading  $S$ 's blocks one by one into the remaining buffer emptying and saving each bucket to memory as it is filled, Relation  $R$  is partitioned in the same way, the hash function  $h$  ensures that tuples with the same value for the join attribute end up in the  $i_{th}$  bucket in the two relations.

**Join/Probe Phase** Going from 1 to  $M - 1$ ,  $\min(R_i, S_i)$  according to number of blocks where  $X_i$  is the  $i_{th}$  partition of  $X$  is loaded to memory and block by block of  $S_i$  is loaded to memory where every tuple  $s$  of  $S_i$  is compared with every tuple  $r$  of  $R_i$  and  $r \circ s$  is outputted whenever  $s.A = r.B$  where  $A$  and  $B$  are  $S$  and  $R$ 's join attribute respectively.

---

Hashing  $S$  and  $R$ 


---

```

foreach tuple  $s \in S$  do
  put  $s$  in bucket  $h(s.A)$ 
end
write buckets to secondary memory
foreach tuple  $r \in R$  do
  put  $r$  in bucket  $h(r.B)$ 
end
write buckets to secondary memory

```

---



---

Probing

---

```

for  $k = 1$  to  $M - 1$  do
   $X \leftarrow \min(S_k, R_k)$ 
   $Y \leftarrow \max(S_k, R_k)$ 
  foreach tuple  $x \in X$  do
    foreach tuple  $y \in Y$  do
      if  $x.A = y.B$  then
        output  $x \circ y$ 
      end
    end
  end
end

```

---

## Nested Loop Joins

The simplest join method where  $S$  is loaded into memory  $M - 1$  blocks at a time in  $\lceil \frac{B(S)}{M-1} \rceil$  iterations where  $B(S)$  is the number of blocks of  $S$  and  $R$  is loaded block at a time in the remaining buffer so that each tuple of  $S$  is compared to each tuple of  $R$  and the tuples where the join attribute is equal are outputed.

---

Nested Loop Join Algorithm

---

```

foreach  $M - 1$  blocks of  $S$  do
  foreach block  $b$  of  $R$  do
    foreach tuple  $t$  in  $b$  do
      foreach tuple  $s$  in the  $M-1$  blocks loaded do
        if  $t.B=s.A$  then
          output  $r \circ s$ 
        end
      end
    end
  end
end

```

---

## Merge Joins

The Merge Join takes two relations and sorts them if not already sorted. It loads one block of each relation and takes the minimum first element according to the join attribute with value  $m$ , then checks whether  $m$  is present at the top of the second relation. If not, tuples with  $m$  are dropped and the process repeats. If  $m$  is indeed found in the second relation, all tuples having  $m$  in the second relation are joined with the initial tuple loading as many blocks as necessary. The process repeats until one relation is exhausted.

---

Merge Join Algorithm

---

```

set  $i \leftarrow 1, j \leftarrow 1$ 
while  $(i \leq n)$  and  $(j \leq m)$  do
  if  $R[i].A > S[j].B$  then
    set  $j \leftarrow j + 1$ 
  end
  else if  $R[i].A < S[j].B$  then
    set  $i \leftarrow i + 1$ 
  end
  else
    output  $R[i] \circ S[j]$ 
    set  $z \leftarrow j + 1$ 
    while  $(z \leq m)$  and  $(R[i].A = S[z].B)$  do
      output  $R[i] \circ S[z]$ 
      set  $z \leftarrow z + 1$ 
    end
    set  $k \leftarrow i + 1$ 
    while  $(k \leq n)$  and  $(R[k].A = S[j].B)$  do
      output  $R[k] \circ S[j]$ 
      set  $k \leftarrow k + 1$ 
    end
    set  $i \leftarrow k, j \leftarrow z$ 
  end
end

```

---

## Question 2

### Bitmap Indices

A Bitmap index is constructed for a single value  $v$  of a single column in a relation  $R$  where  $R$ 's tuples are numbered from 0 to  $n$ . The index for  $v$  is a vector of  $n$  bits where if the  $n_{th}$  tuple of  $R$  has value  $v$  then the  $n_{th}$  bit of the vector is a 1 and 0 otherwise, as an example consider the relation Student.

Name	Age	Hobby
Osama	21	swimming
Ahmed	29	hiking
Ayman	20	reading
Bassem	22	swimming
Yasser	22	reading
Mostafa	21	reading

indices for Age :

20 : 001000

21 : 100001

22 : 000110

indices for Hobby :

swimming : 100100

reading : 001011

hiking : 010000

to add efficiency to Bitmaps run length encoding and other compression techniques can be used to minimize the number of bits as much as possible, for deletion and insertion properties the number associated with tuples is considered to be constant so that when inserting a new tuple in a relation simply a new bit is added to all its bitmap indices and when a tuple is deleted a tombstone is inserted in its position in all its bitmap indices, another way of handling indices is by using *existence bitmap* [?] which is a bitmap similar to the others but its  $n_{th}$  bit is 0 when the  $n_{th}$  row is deleted.

Bitmap indices are best used when

1. Insertion/Deletion rate is low.
2. When number of distinct values of a column are low "If the number of distinct values of a column is less than 1% of the number of rows in the table, or if the values in a column are repeated more than 100 times, then the column is a candidate for a bitmap index."